

CSE 360-Computer Architecture

Lecture-4


Cache Memory

Dr. Shamim Akhter

Computer Science and Engineering



Processor-Memory Performance Gap



Large and Fast: Exploiting Memory Hierarchy

Illusion of unlimited fast memory

Introduction




Unlimited memory




Google Data Center

Memory Facts



How to create such memory



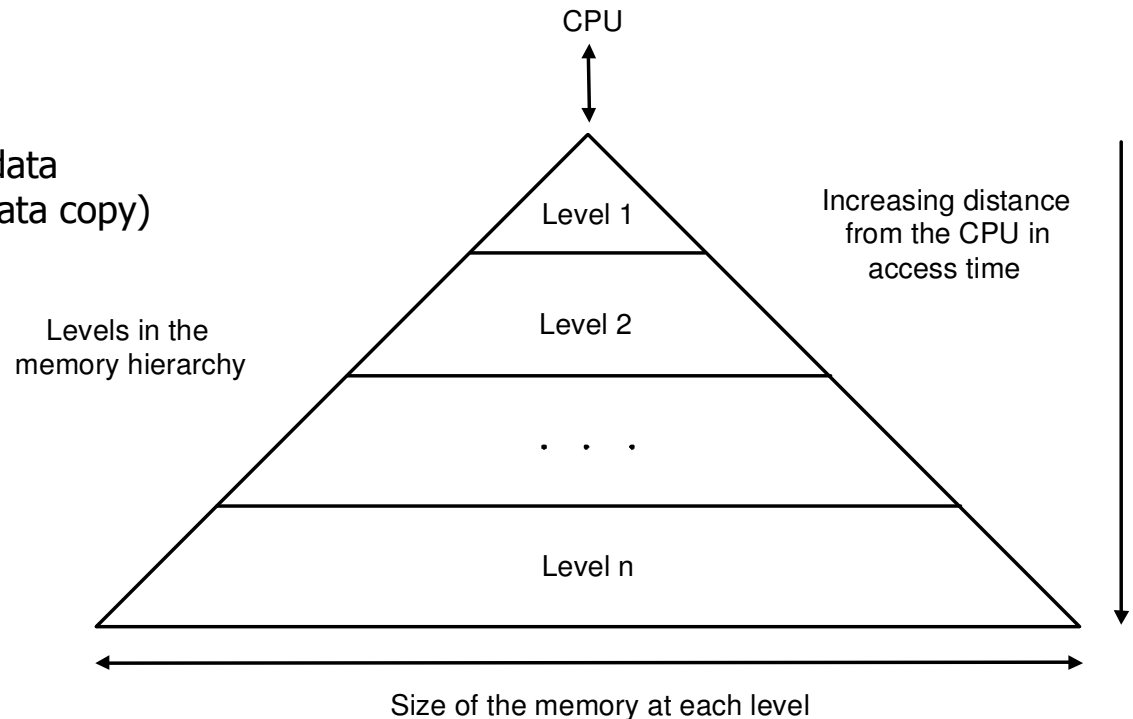
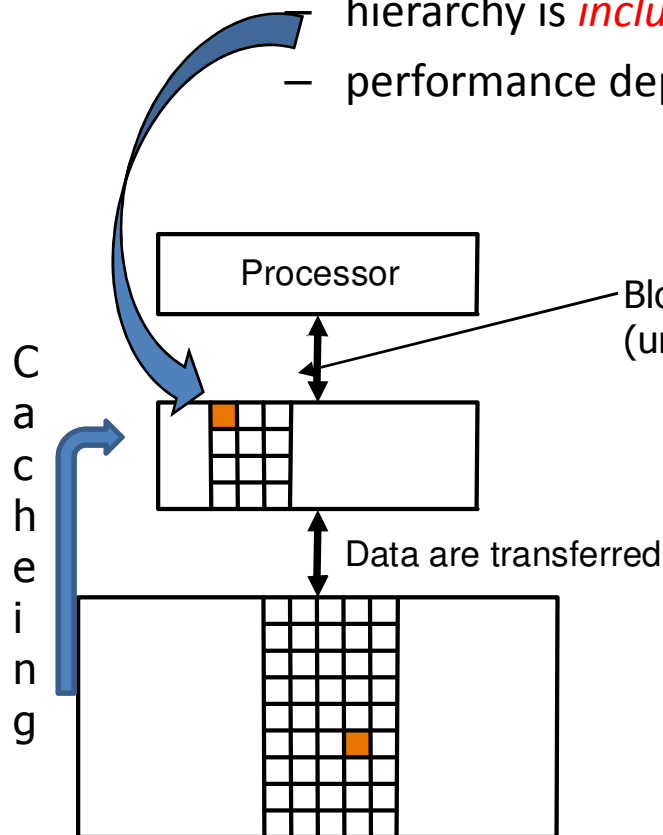
Illusion of being large, cheap and fast



Memory Hierarchy

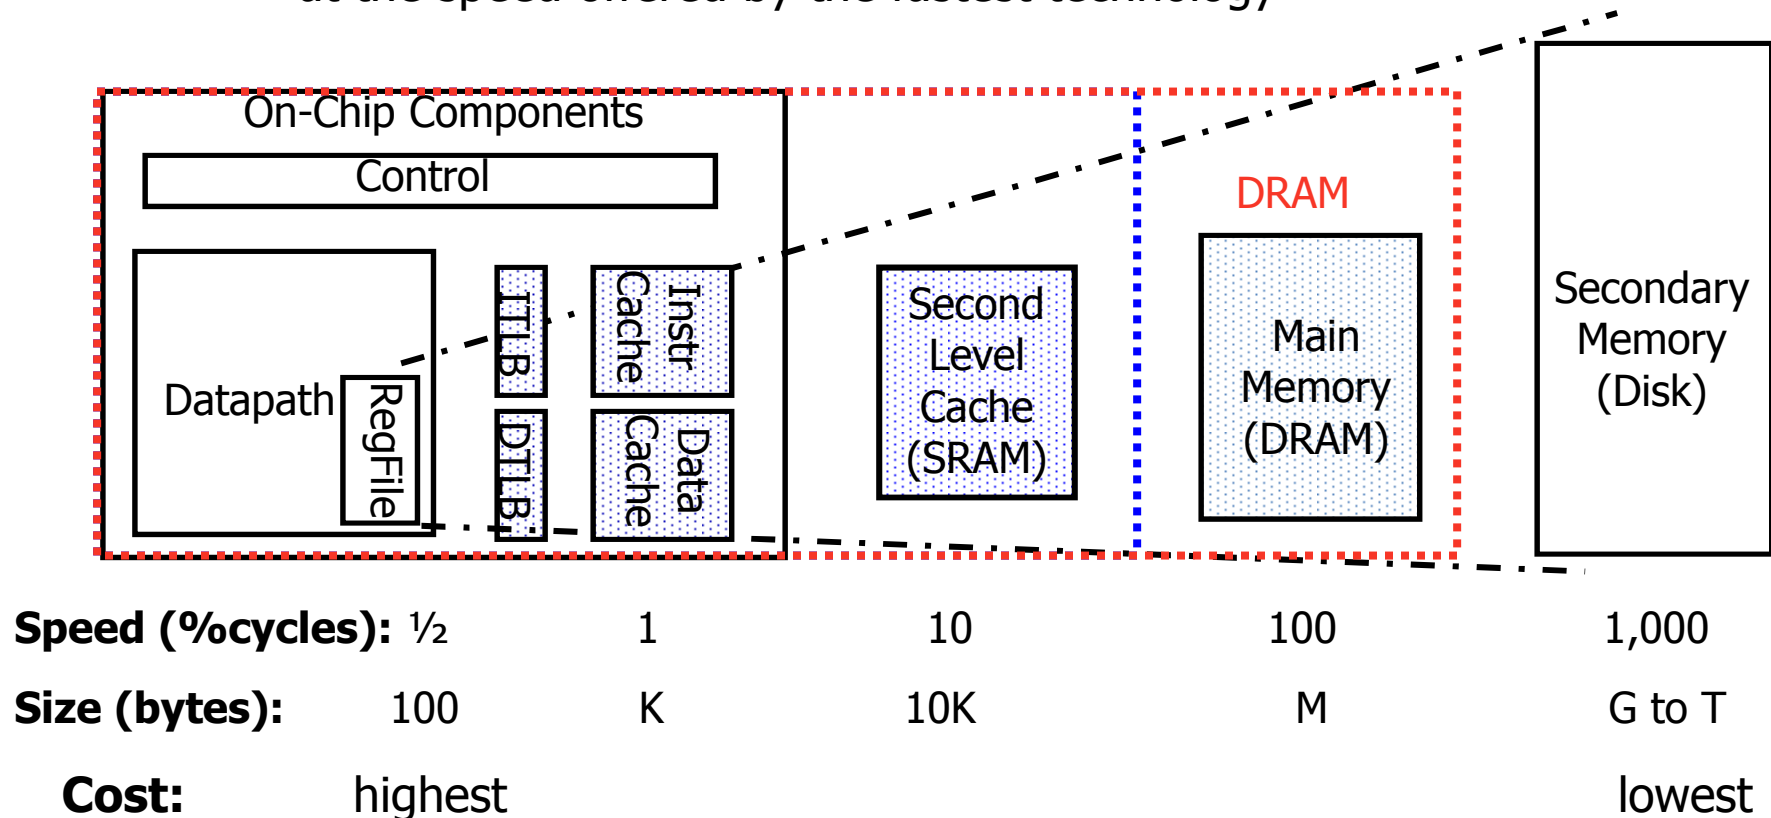
Memory Hierarchy

- Users want large and fast memories...
 - expensive but they don't like to pay...
- Make it seem like (illusion)- **they have what they want...**
 - *memory hierarchy*
 - hierarchy is **inclusive**, every level is subset of lower level
 - performance depends on *hit rates*



A Typical Memory Hierarchy

- By taking advantage of memory hierarchy
 - Can present the user
 - with as much memory as is available
 - in the cheapest technology
 - at the speed offered by the fastest technology



Cache Memory

A buffer memory that holds the currently most useful program segments and data, in much the same way that registers hold some of the currently active data elements for rapid access by the processor.

@ Behrooz Parhami

How does cache improve system performance?

Locality Principle

Cache applied locality principle to minimize the calls to memory.

Temporal locality

Same item will tend to be referenced soon

Spatial locality

Once an instruction/data has been accessed, it /or its neighbors will be referenced soon.

Sequentiality

Reference made to a particular location s then it is likely to reference $s+1$ within the next several references. Sequentiality is a restricted type of spatial locality.

```
for(i = 0; i < 20; i++)  
  for(j = 0; j < 10; j++)  
    a[i] = a[i]*j;
```

Outer loop is an example of spatial locality.

-It sequentially increments the address of the inner for- loop calls.

A[0]	A[1]	A[2]	A[3]	A[19]
------	------	------	------	------	-----	-------


Inside loop demonstrates temporal locality.

▪ Same memory address is accessed ten times in a row, and multiplied by j each time.

▪ Both i and j (loop counters) are very good examples of temporal locality.

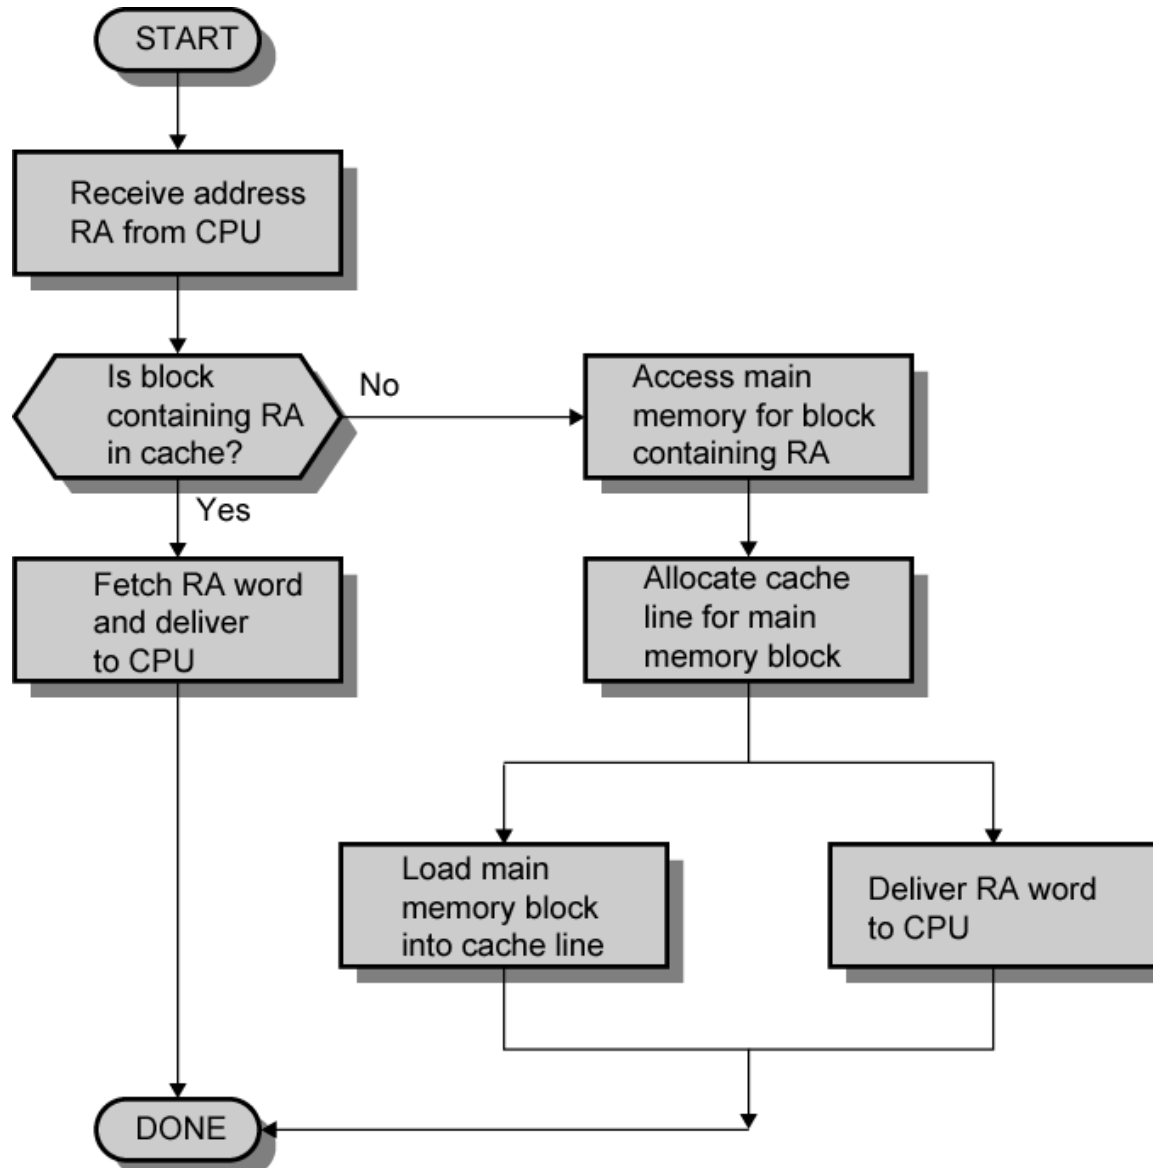
A[0]*0	A[0]*1	A[0]*2	A[0]*3		A[0]*9
--------	--------	--------	--------	------	--	--------

Cache and Main Memory Architecture

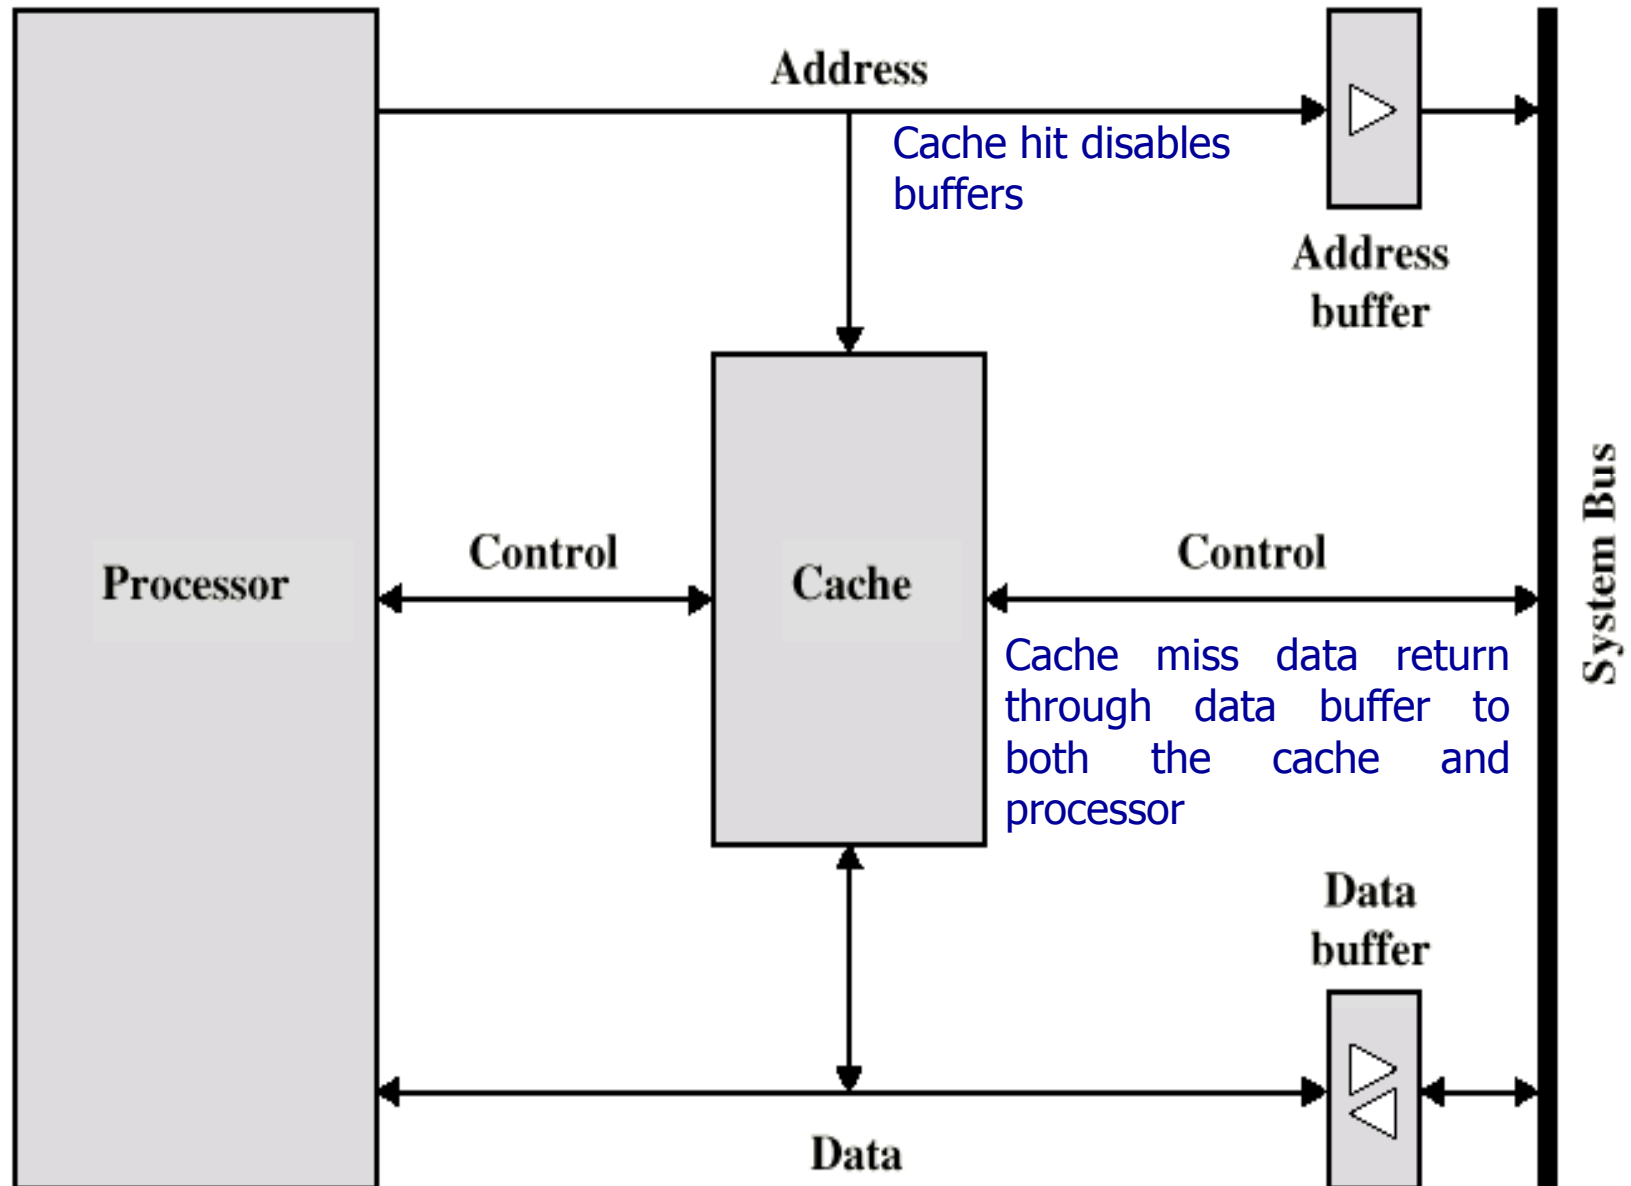


Cache Read Operation

RA=Read
Address




Cache Organization



Caches

- By simple example
 - assume block size = one word of data = 32 bits






Issues:

- I. how do we know if a data item is in the cache?
- II. if it is, how do we find it?
- III. if not, what do we do?

Solution depends on *cache addressing scheme*...

Cache Addressing Schemes



Direct- Mapped 32 bytes Cache

Tag	Index	Byte Offset
3	3	2

	v	TAG	Block Data			
0						
1	1		33	36	39	21
2						
3						
4						
5						
6						
7						



1 Block = 4 bytes

Total Cache Block = $32/4 = 8$ Blocks

Block Address

0	11	12	13	11	0-3
1	12	13	1	5	4-7
2	7	9	2	11	8-11
3	2	4	3	12	12-15
4	3	5	1	11	16-19
5	6	7	9	1	20-23
6	9	8	11	0	24-27
7	11	13	17	20	28-31
8	23	12	22	27	32-35
9	33	36	39	21	36-39
10	47	45	42	43	40-43
11	101	212	231	22	44-47
12	11	134	22	33	48-51
13	11	22	43	44	52-55
14	223	34	33	0	56-59
15	8	9	67	54	60-63
16	4	4	4	4	64-67
17	65	53	33	22	68-71
18	88	33	54	25	72-75
19	90	96	11	34	76-79
20	49	43	22	43	80-83
21	90	67	22	21	84-87
22	53	78	87	21	88-91
23	44	55	23	33	92-95
24	55	2	7	55	96-99

Byte Address

Questions: Direct Mapped Cache


1. How do you define the main block and its corresponding cache block?
2. How to check the data availability in cache or not?

Solution 1: Simple

$\text{Cache block\#} = \text{main block\#} \% \text{Total Cache block}$


Direct Mapped Cache

- Addressing scheme in *direct mapped* cache:
 - cache block address = memory block address % cache size
(*unique*)
 - if cache size = 2^m ,
 - cache address = lower m bits of n -bit memory address
 - remaining upper $n-m$ bits kept as *tag bits* at each cache block
 - also need a *valid bit* to recognize valid entry



How to check the data availability in cache or not?


- Need to decide the address formats
 - 256(2^8)bytes main memory
 - 4 bytes block size
 - # of blocks $256/4 = 64$
- Cache size 8 blocks
- Lets ref. **39 address** from main memory
 - Bit address 00100111



Direct Mapped Cache

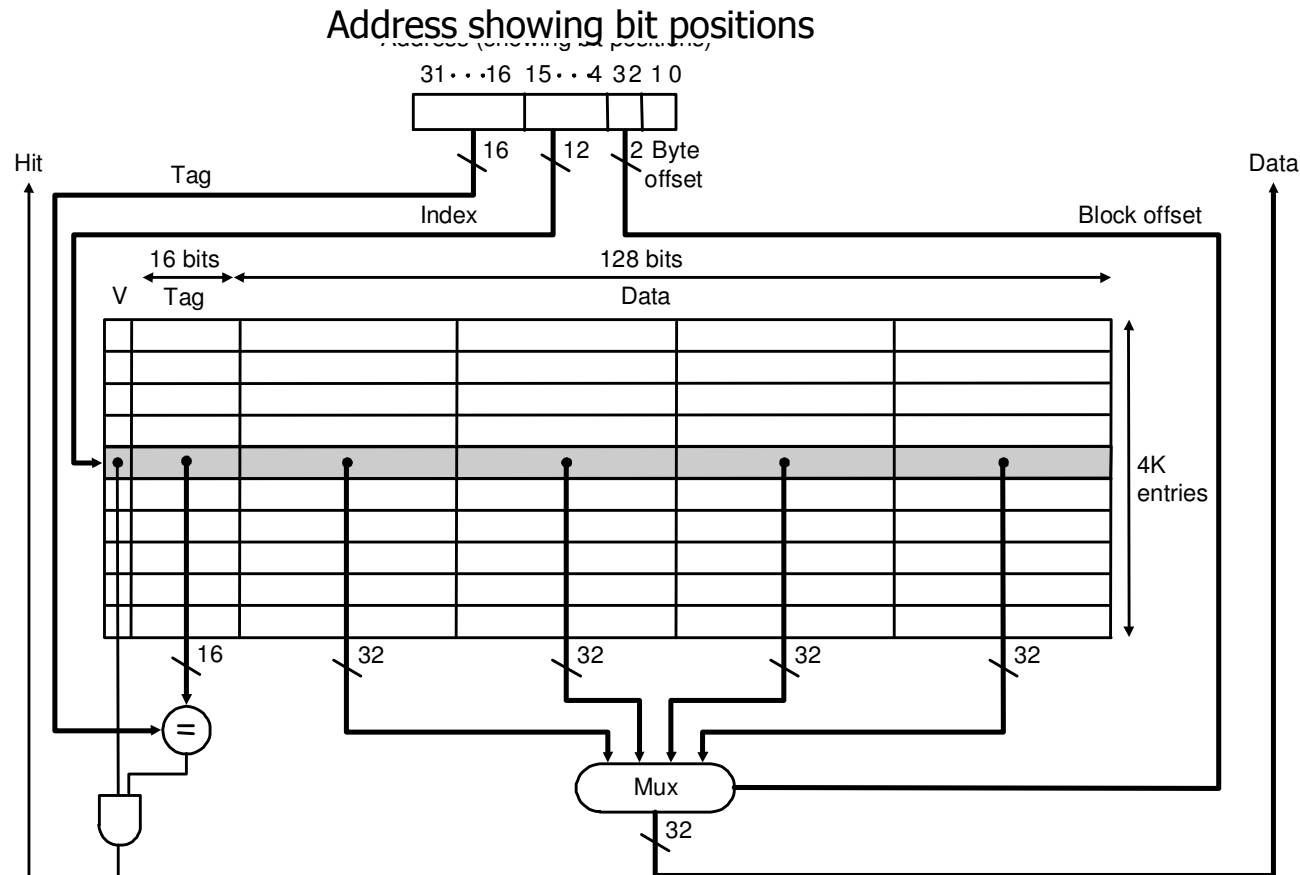
- MIPS style:

Recently address words replace less recently referenced words.



Direct Mapped Cache: Taking Advantage of Spatial Locality

- Taking advantage of spatial locality with *larger* blocks:



Cache with 4K 4-word blocks: *byte offset* (least 2 significant bits) is ignored, next 2 bits are *block offset*, and the next 12 bits are used to index into cache

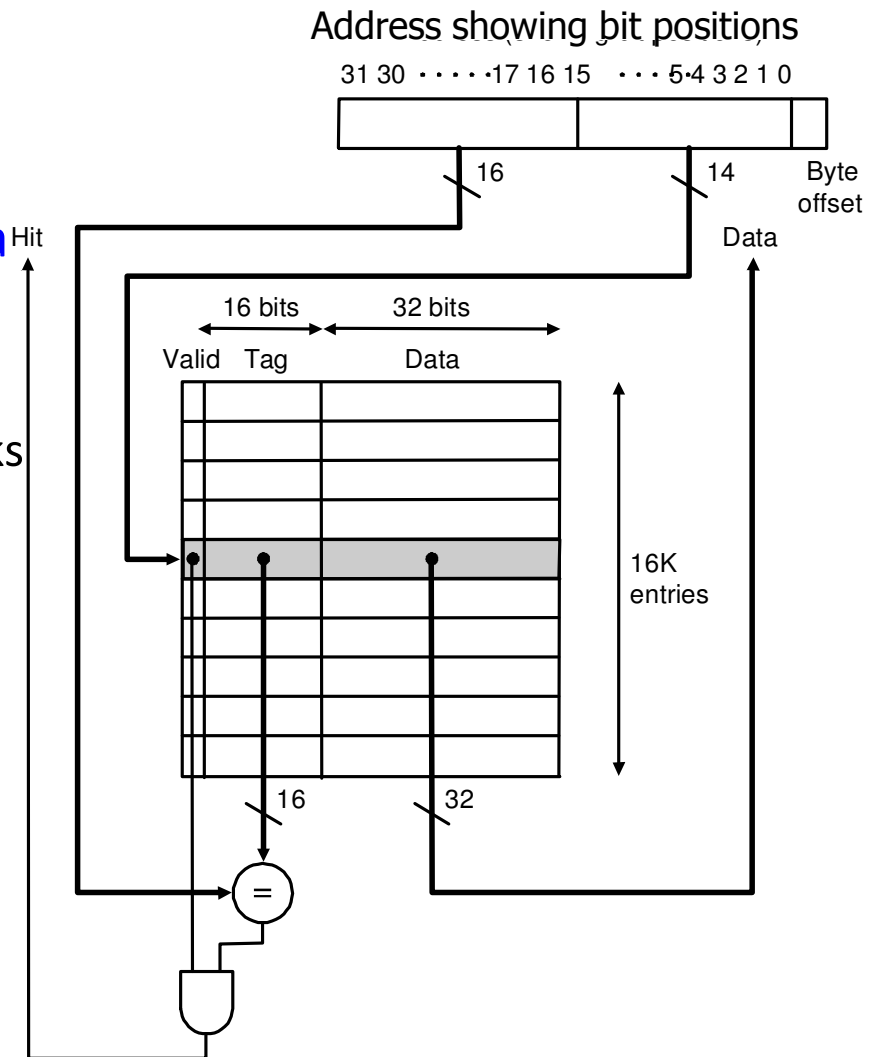
Example Problem

- In a Direct Mapped Cache:
- Consider a cache with **64 blocks and a block size of 16 bytes**. What block # does the byte address **1200** map to?
- The block is given by
$$(\text{Block address}) \bmod (\# \text{ of cache blocks})$$
- Where the block address is
$$\text{Byte address} / \text{Bytes per block}$$
- Thus, with 16 bytes per block, byte address 1200 is block address
$$\text{Floor}(1200/16) = 75$$
- which maps to cache block # **$(75 \bmod 64) = 11$**

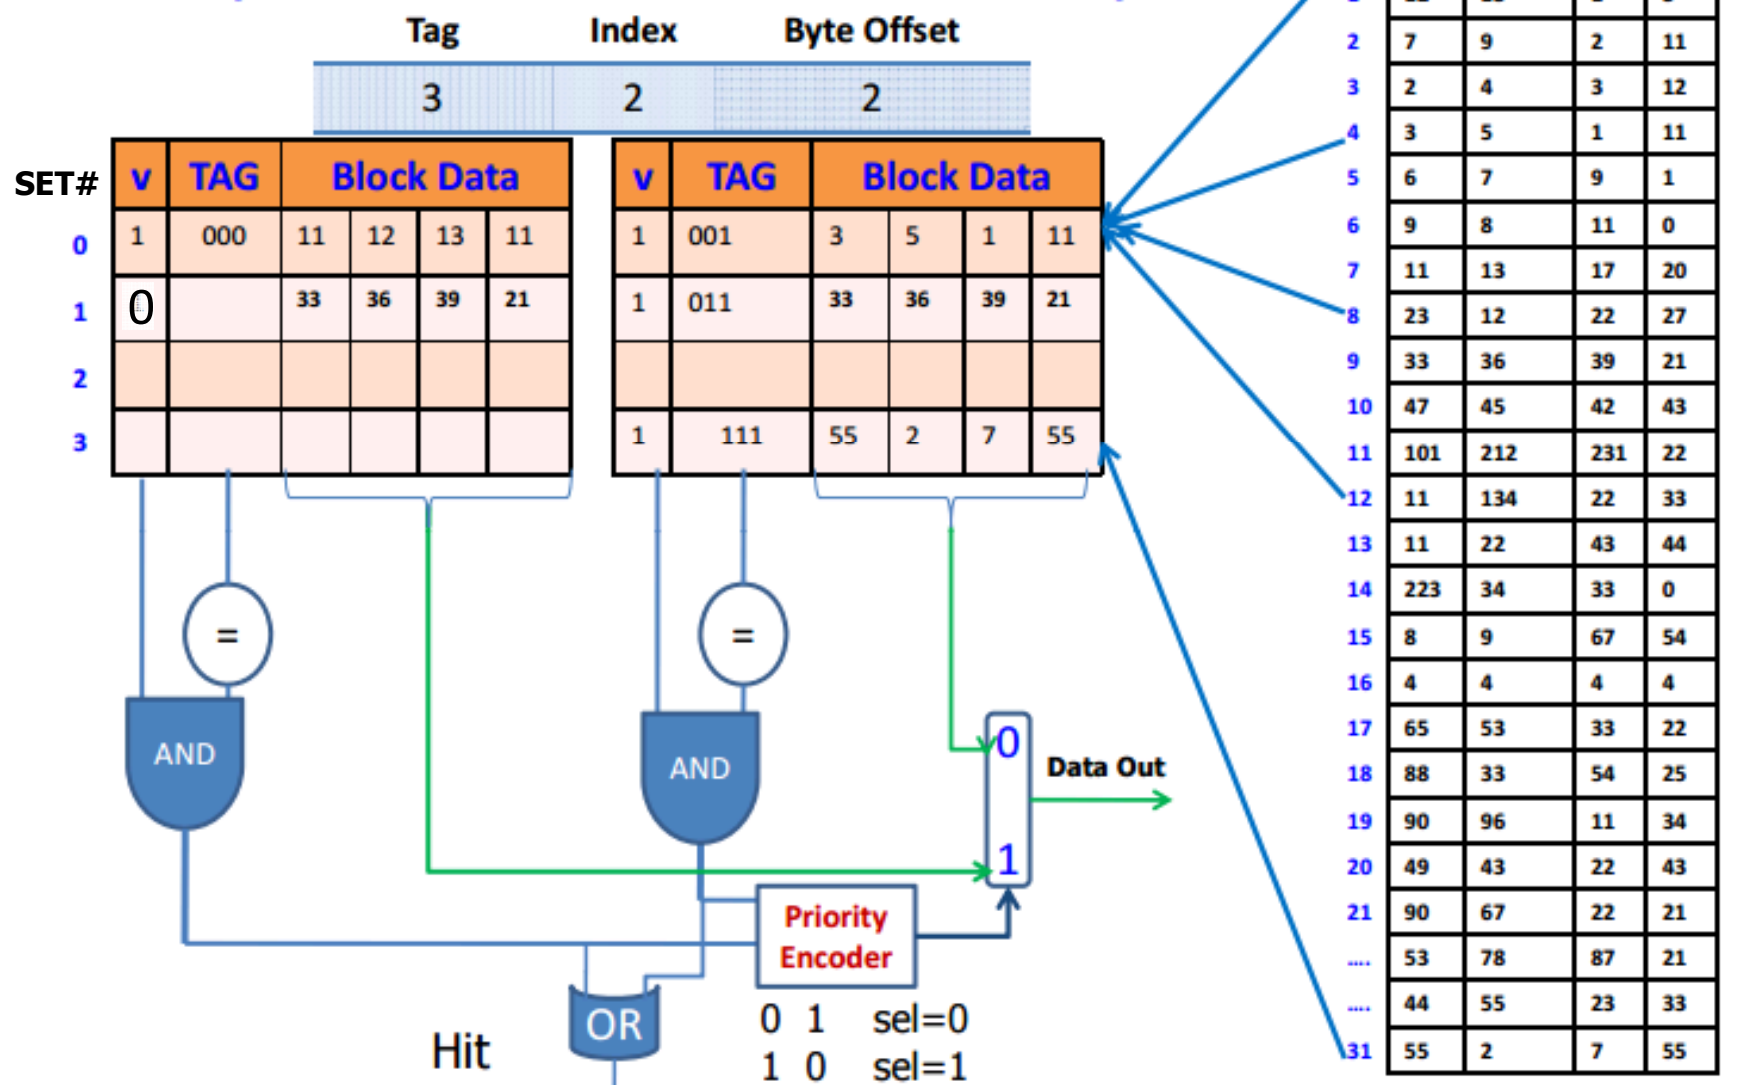
Example: DECStation 3100 Cache (MIPS R2000 processor)

Cache with 16K 1-word blocks: *byte offset* (least 2 significant bits) is ignored and next 14 bits used to index into cache

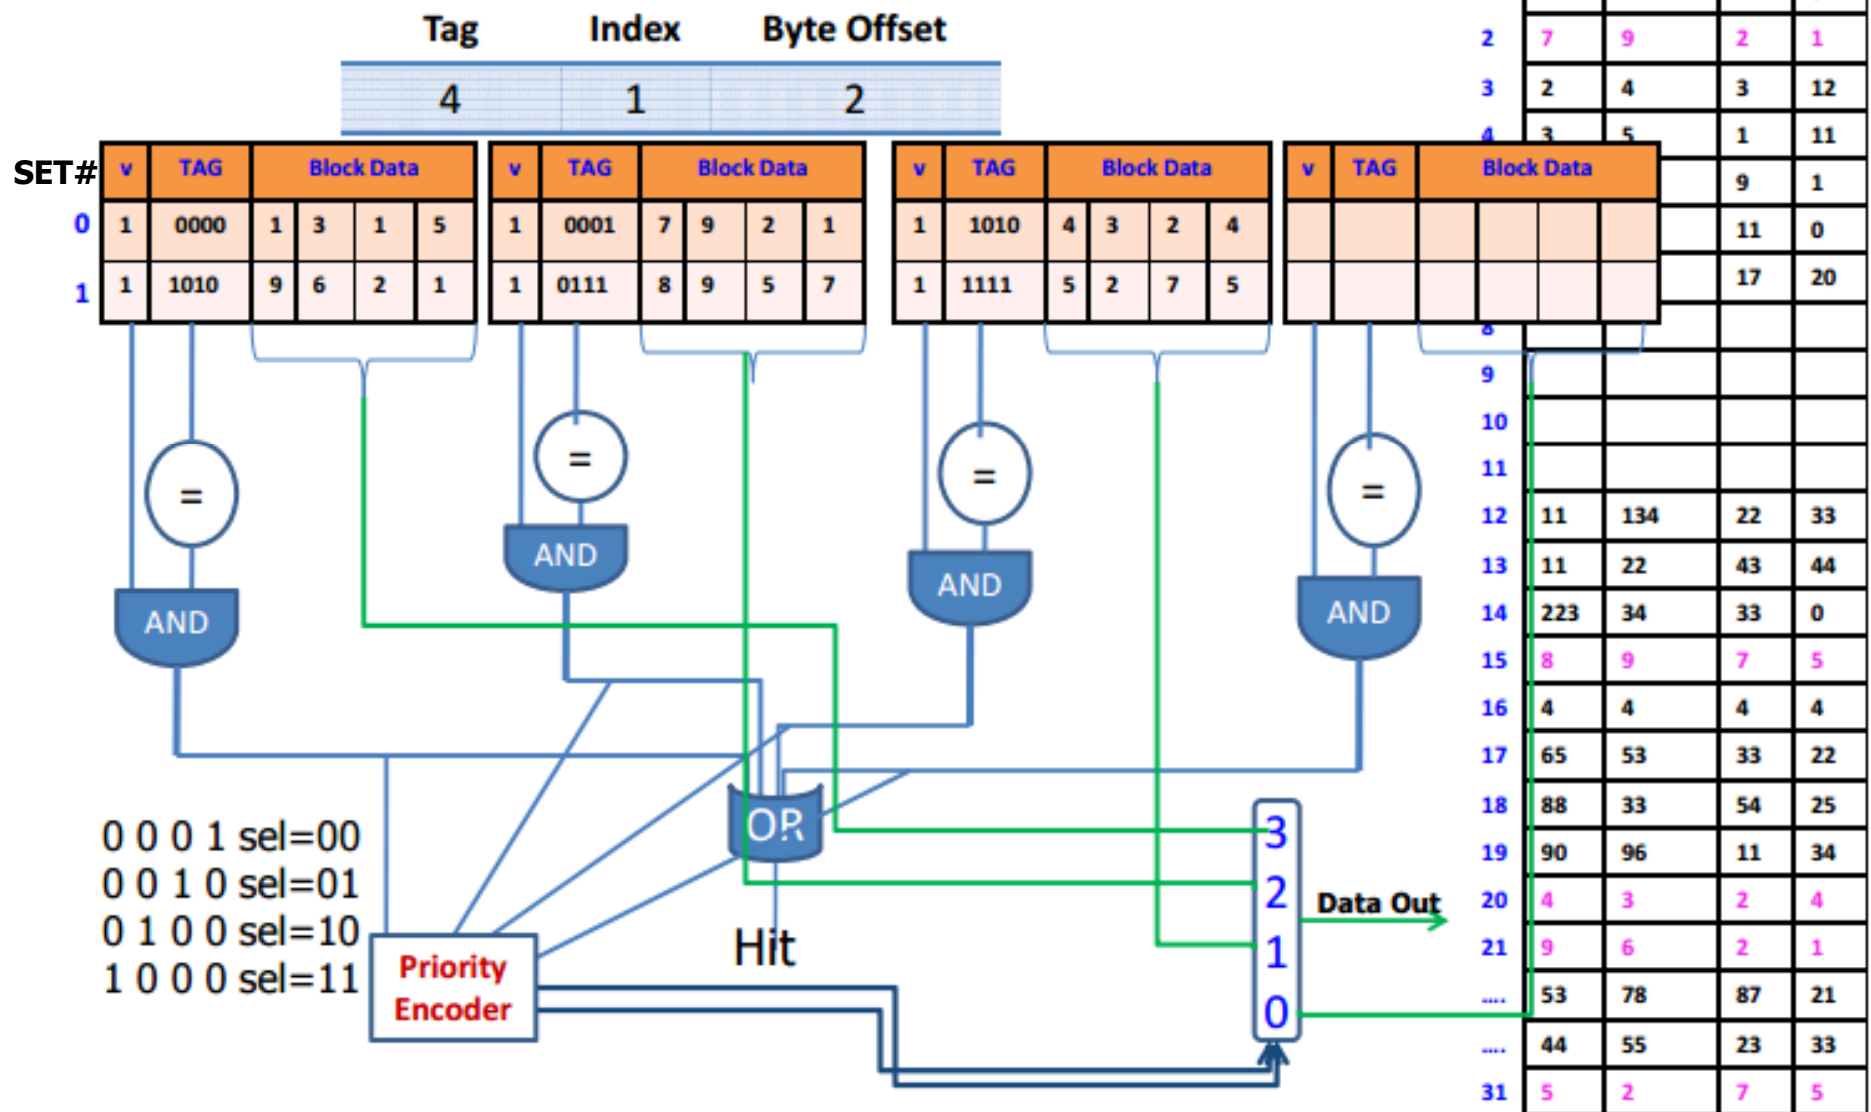
- How many total bits are required for a direct-mapped cache with 64 KB of data in total and one-word blocks, assuming a 32-bit address?
- $64 \text{ KB} = 16\text{K Word} = 16 \times 1024 = 2^{14} \text{ words} = 2^{14} \text{ blocks}$
- Thus, $\text{Tag} = (32 - 14 - 2) = 16 \text{ bits}$
- Each 2^{14} words has 16bits tag+ 1 bit validity bit
- Total Size = $2^{14} \times (32 + 17) = 2^{10} \times 784 = 784 \text{ Kbits} = 98\text{KB}$
- Cache size ~1.5 times bigger than total storage data



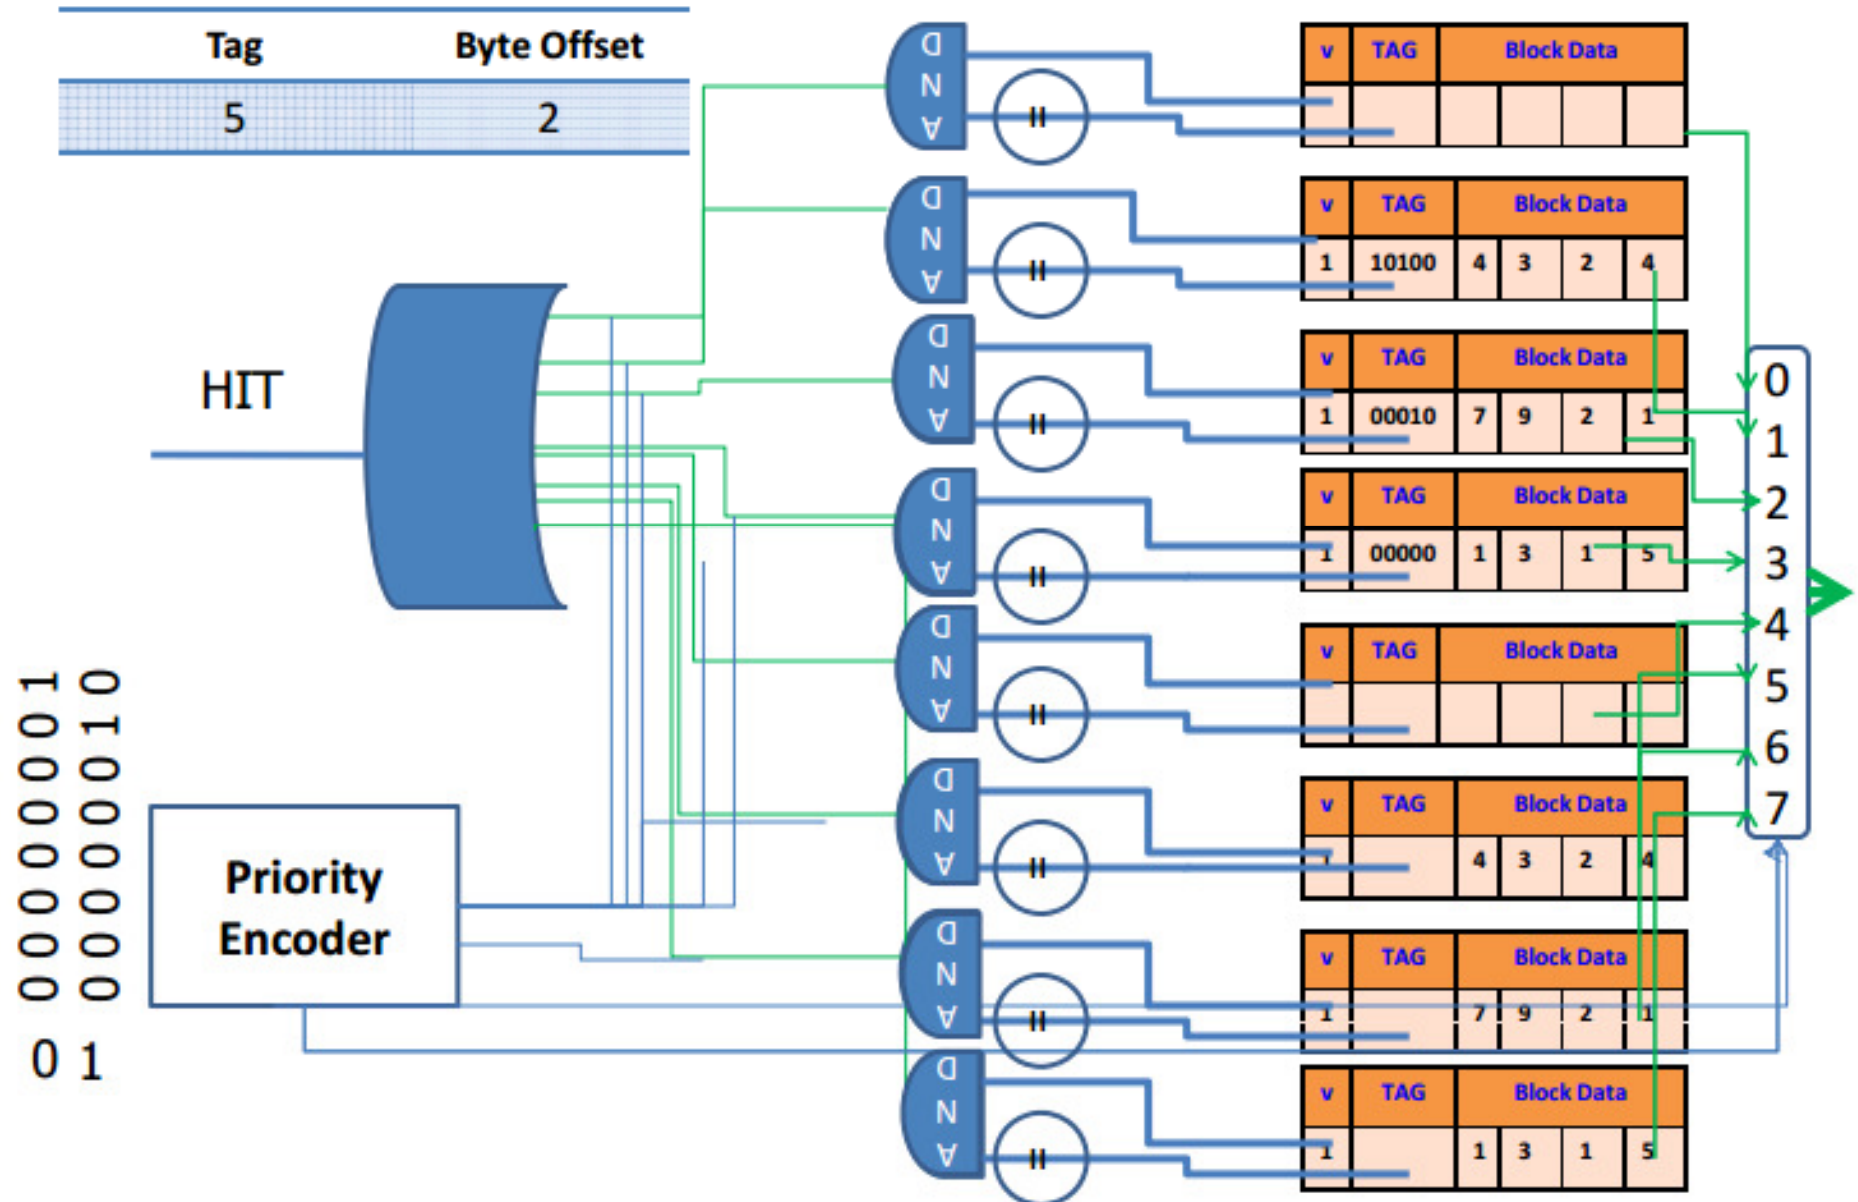
2-Way Set Associative- 2 blocks per set 32 bytes Cache- Block size 4 bytes



4-Way Set Associative- 4 blocks per set 32 bytes Cache- Block size 4 bytes



Fully Set Associative- All Blocks in One Set



- If the total size is kept the same,
 - increasing associativity increases the # of blocks per set (# of comparators)
 - Increase factor of 2 in associativity
 - doubles the # of blocks per set and halves the # of sets.
 - decreases the size of the index by 1 bit
 - increases the size of the tag by 1 bit.
 - Fully associative cache
 - Only one set and all blocks must be checked in parallel
 - No index
 - Tag address address-block offsets
 - Direct-mapped cache
 - Single comparator is needed (only one block)

What Block Should be Replaced on a Cache READ Miss?

- When a miss occurs in an **associative cache**- which block to replace?
 - **Random**: Candidate blocks are randomly selected.
 - **Least recently used (LRU)**: Candidate block is the one that has been unused for the longest time.
 - **First in First Out (FIFO)**: Selects the oldest rather than the LRU block.

Example Problems

- Find the number of misses for a cache with four 1-word blocks given the following sequence of memory block accesses:
0, 8, 0, 6, 8
for each of the following cache configurations
 1. **direct mapped**
 2. **2-way set associative (use LRU replacement policy)**
 3. **fully associative**

Solution

- 1 (direct-mapped)

Block address	Cache block
0	0 ($= 0 \bmod 4$)
6	2 ($= 6 \bmod 4$)
8	0 ($= 8 \bmod 4$)

Block address translation in direct-mapped cache

Address of memory block accessed	Hit or miss	Contents of cache blocks after reference			
		0	1	2	3
0	miss	Memory[0]			
8	miss	Memory[8]			
0	miss	Memory[0]			
6	miss	Memory[0]		Memory[6]	
8	miss	Memory[8]		Memory[6]	

- 5 misses

Solution (cont.)

- 2 (two-way set-associative)

Block address	Cache set
0	0 ($= 0 \bmod 2$) mod set number
6	0 ($= 6 \bmod 2$)
8	0 ($= 8 \bmod 2$)

Block address translation in a two-way set-associative cache

Address of memory block accessed	Hit or miss	Contents of cache blocks after reference			
		Set 0	Set 0	Set 1	Set 1
0	miss	Memory[0]			
8	miss	Memory[0]	Memory[8]		
0	hit	Memory[0]	Memory[8]		
6	miss	Memory[0]	Memory[6]		
8	miss	Memory[8]	Memory[6]		

- Four misses

LRU page replacement

Solution (cont.)

- 3 (fully associative)

Address of memory block accessed	Hit or miss	Contents of cache blocks after reference			
		Block 0	Block 1	Block 2	Block 3
0	miss	Memory[0]			
8	miss	Memory[0]	Memory[8]		
0	hit	Memory[0]	Memory[8]		
6	miss	Memory[0]	Memory[8]	Memory[6]	
8	hit	Memory[0]	Memory[8]	Memory[6]	

- 3 misses

What Happens on a Write?

- Two approaches
 - Write-Through Scheme
 - Write-back Scheme

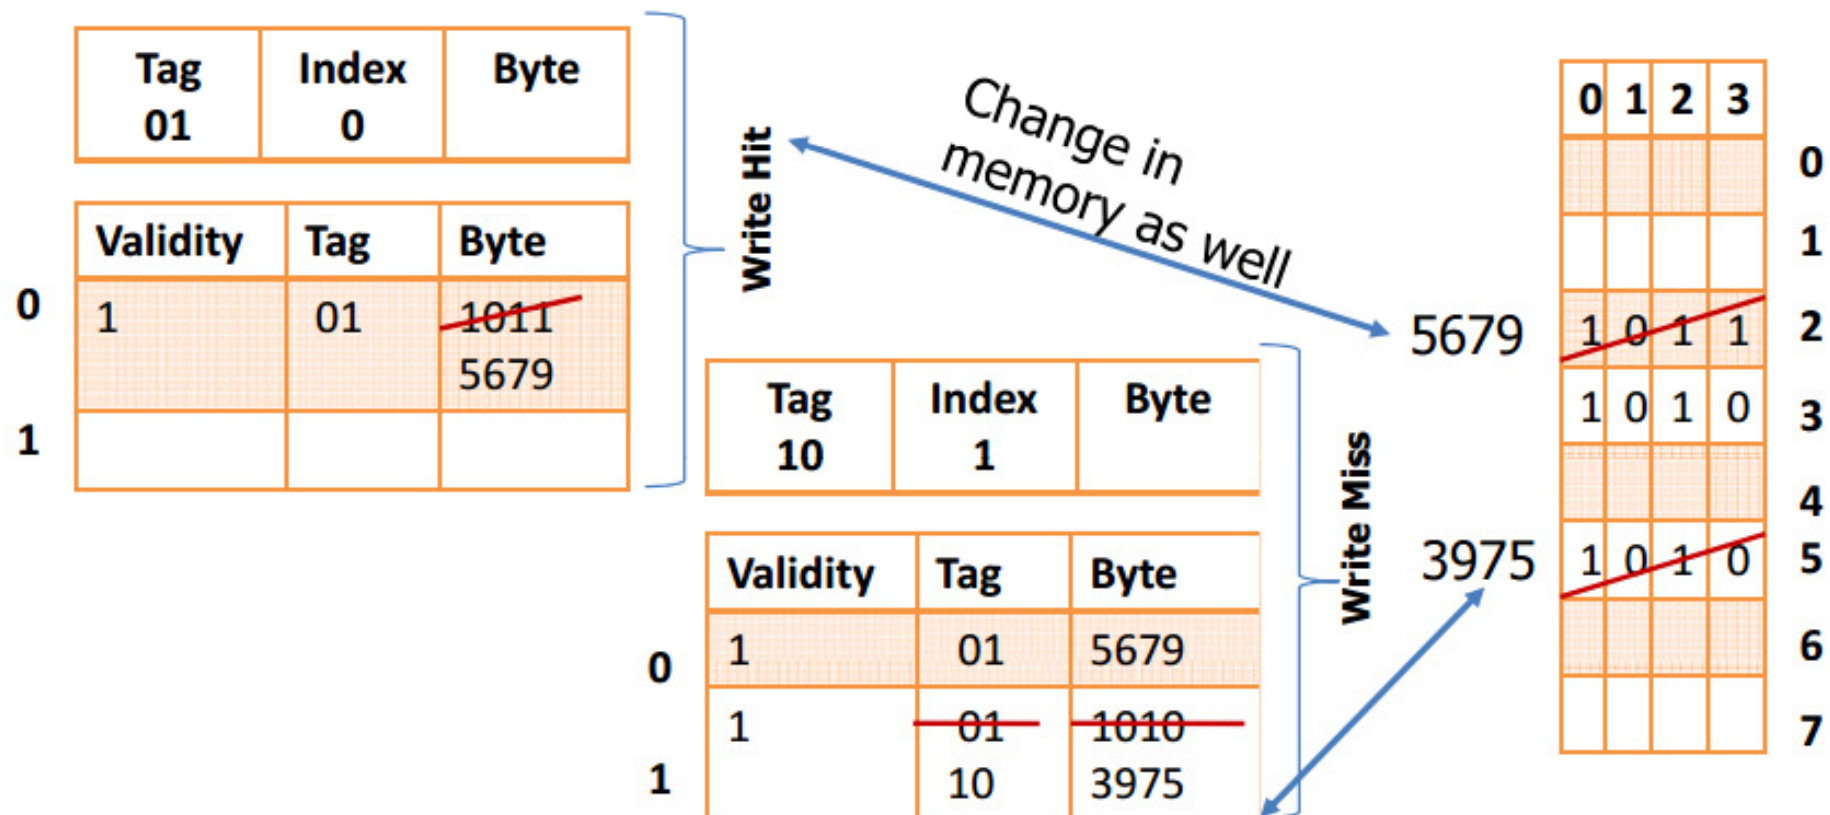
Cache Write Hit/Miss (Ex. store op)

1-word blocks

- Write-through scheme**

- on **write hit**: replace data in cache and memory blocks with every write hit to avoid inconsistency
- on **write miss**: update tag, valid bit, write the word into cache and memory


X obviously no need to read missed word from memory!



Disadvantage of Write-through Memory Write

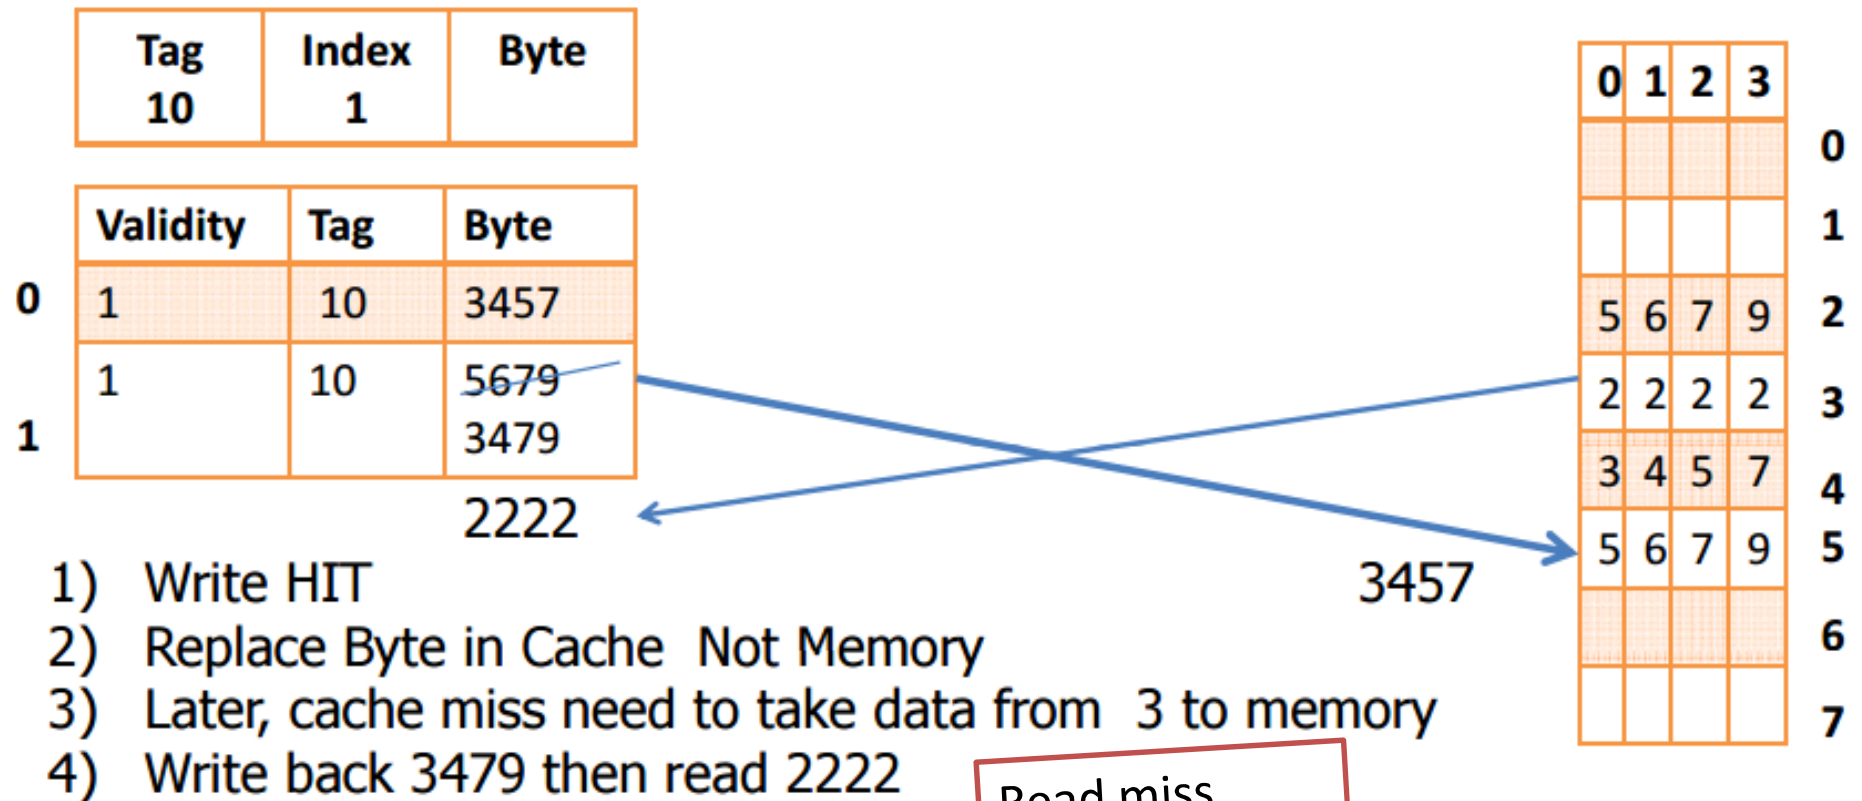
– Write-through is slow because of always required memory write

- performance is improved with a **write buffer** where words are stored while waiting to be written to memory – processor can continue execution until write buffer is full.
- when a word in the write buffer completes writing into memory that buffer slot is freed and becomes available for future writes
- DEC 3100 write buffer has 4 words



• Write-back scheme

- write the data block only into the cache and write-back the block to memory only **when it is replaced in cache**
- more efficient than write-through, more complex to implement



- 1) Write HIT
- 2) Replace Byte in Cache Not Memory
- 3) Later, cache miss need to take data from 3 to memory
- 4) Write back 3479 then read 2222

Read miss

Improving Cache Performance

- Use split caches for instruction and data because there is more spatial locality in instruction references:


Program	Block size in words	Instruction miss rate	Data miss rate	Effective combined miss rate
gcc	1	6.1%	2.1%	5.4%
	4	2.0%	1.7%	1.9%
spice	1	1.2%	1.3%	1.2%
	4	0.3%	0.6%	0.4%

**Miss rates for gcc and spice in a MIPS R2000
with one and four word block sizes**

Direct Mapped Cache

Taking Advantage of Spatial Locality

- Miss rate falls at first with increasing block size as expected, but, as block size becomes a large fraction of total cache size, miss rate may go up because
 - there are few blocks
 - competition for blocks increases
 - blocks get ejected before most of their words are accessed (*thrashing* in cache)

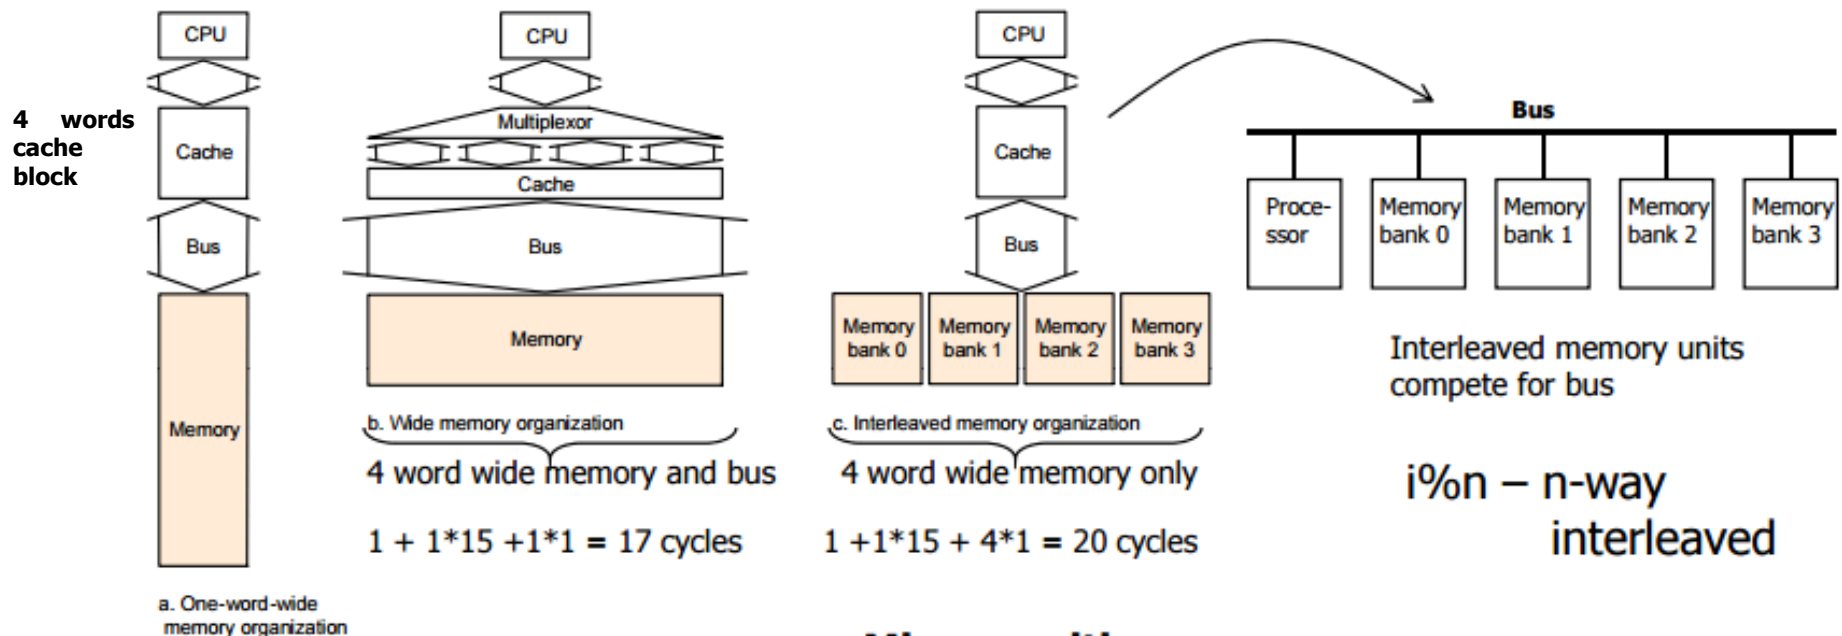


Improving Cache Performance by Increasing Bandwidth

- Assume:

- cache block of 4 words
- 1 clock cycle to send address to memory address buffer (1 bus trip)
- 15 clock cycles for each memory data access
- 1 clock cycle to send data to memory data buffer (1 bus trip)

Interconnectivity



Miss penalties

Memory wide 2 words and bus